

Is Civilization Anti-Nature?

A Conditional Argument Against Capability Elimination, with Capability’s Insufficiency
and a Viability Constraint for Alignment

Ajax Benander*

September 2026 — Working Draft, rev. 8.2

Abstract

Two positions occupy the ends of a spectrum about the human role. Capability-first techno-optimism holds that capability is the thing and more of it is the answer; environmental extinctionism holds that human extinction would serve the living world. This paper argues, under one explicit normative premise and a survivorship model proved in a companion, that both fail, and in structurally opposite ways. The claim in one sentence: for an agent who values the biosphere’s long-run persistence, removing the mechanism that maintains its hazard-muffling capability, so that accumulated capability decays to zero, cannot be an environmental improvement, since the total hazard then has a permanent floor and viability is exactly zero; while maximizing capability is not by itself sufficient for the objective, since viability requires a favorable response sign together with favorable driver dynamics, and no magnitude of capacity substitutes for the sign. Four things are established at their stated grades. The two-by-two picture: capability decaying to zero is nonviable at either sign, capability accumulating is potentially viable only at the favorable sign, so that the model’s live region is accumulating capability with favorable anatomy; elimination forfeits the option, while maximization is insufficient without favorable anatomy and is not thereby shown suboptimal, since for fixed $\rho > 0$ both hazard channels decrease with x , and an optimum elsewhere would need an optimization problem with costs, controls, constraints and anatomy dynamics. The winner’s irony, stated as a proposition rather than a sentiment: in the paper’s homogeneous pinned prerequisite schedule, the viability forfeited by abandonment is nondecreasing in the gates already cleared, while historical surprisal is tracked separately and is also nondecreasing along that schedule. Thus, on geological timescales, the lineage has only just come into possession of the biosphere’s sole realized route to positive viability, and to dismantle it for the biosphere’s sake is, within that model, to push away winnings out of ignorance of what has just happened, with the fate of the living world as the stake. The concession, that a successor muffling agent might re-evolve the route, priced as a competing-risks residual rather than argued away. And a consequence for alignment: the model’s viability boundary is preference-independent, and for any alignment objective constrained to preserve biospheric persistence it is a necessary condition, not the whole objective, and not certifiable from finite observation of a trajectory; treating the boundary as a normative constraint depends on the paper’s one premise and on nothing in the mathematics. The paper owns its premise, names its targets narrowly, and ends on the conditional the

*Email: akb9408@rit.edu. This paper draws the consequences, for three positions about the human role, of a survivorship model proved in a companion [1], a mechanism developed in a second [2], and a target constructed and a hypothesis stated in a third [3]; each is cited and restated, none presupposed beyond what is stated, and the conditions under which each conclusion holds are named with it. Drafting collaboration: Anthropic Claude (Opus 4.x, Fable 5), under the author’s direction.

model establishes: civilizational disappearance is not the restoration of a self-sustaining biosphere but the removal of the realized route by which exogenous hazard can be endogenously reduced.

1 Introduction

The question whether civilization is anti-nature is usually argued from values, and the arguments do not meet. This paper argues it from a model, under one premise it states rather than derives. The premise is the extinctionist’s own: that greater long-run persistence of the biosphere, considered as a temporally extended living system, is environmentally preferable. The model is a survivorship model proved in a companion, in which a biosphere carries a hazard with two channels, an external one that accumulated capability muffles and an internal one the same capability feeds or starves by the sign of a response exponent. Two results of that model, both simple, do the paper’s work. If capability is pinned at zero, the external channel is unshielded and the internal channel sits at its bare coefficient, so the total hazard has a permanent floor and the cumulative hazard diverges: viability is exactly zero, for every parameter value. And if capability accumulates, viability is the indicator of a favorable sign times a driver-viability functional: it can be positive, and it is never positive at an adverse sign however large the capability. The first result is the argument against environmental extinctionism, on its own valuation. The second is the argument against capability-first techno-optimism. Together they give a two-by-two picture in which eliminating capability forfeits a positive persistence option and maximizing it is insufficient without favorable anatomy, the live region being accumulating capability with favorable anatomy; that maximization is suboptimal is not established, and “persistence” in the normative argument means indefinite-survival probability unless a passage names another criterion, expected duration, finite-horizon survival, welfare or ecological preservation, which are distinct measures, two policies with zero probability of surviving forever differing enormously in expected duration, finite-horizon survival, welfare and ecological preservation. The comparison is strict exactly where maintaining capability preserves a positive continuation prospect that elimination removes; where both the maintained and the eliminated cases have zero indefinite-survival probability in a regime, this criterion alone does not rank them.

A second theme runs through the paper and is stated as a proposition because it is easy to mistake for a sentiment. The lineage that carries the biosphere’s only realized route to the target has, on geological timescales, only just cleared the gates that make the target reachable. In the homogeneous pinned schedule used for the proposition, the viability forfeited by abandoning the route is nondecreasing in the gates already cleared, while the surprisal of having cleared them is also nondecreasing. These are distinct quantities: the retained option is the forward-looking stake; historical surprisal records how unlikely it was to reach the stage and does not itself create value. To wish the route dismantled for the biosphere’s sake is therefore, within that schedule model, to hold a winning ticket and push the winnings away out of ignorance of what has just happened, with the difference that the stake is the continuation of the living world. That is the winner’s irony, and Section 8 gives it its exact form.

Two things the paper does not do are stated at the outset. It does not establish that every alternative lineage is an Imperfection Magnitude Drop; that is a conditional corollary of the Route World companion, under a landscape hypothesis stated there, and it is imported at that grade and used only where marked. And it does not answer the alignment problem’s question of what an artificial agent should be aligned to, which the literature has treated as central for

a decade; it identifies a preference-independent viability boundary, which becomes a constraint that an aligned system must not violate exactly when the alignment objective is constrained to preserve biospheric persistence, and it says plainly that survival ranks none of the things two viable policies can differ in.

The claim, in one sentence. For an agent who values the biosphere’s long-run persistence, removing the mechanism that maintains its hazard-muffling capability, so that the capability decays to zero, cannot be an environmental improvement, since a decayed capability leaves a permanent hazard floor and zero viability; maximizing capability is not by itself sufficient for the objective, since no magnitude of capacity substitutes for a favorable response sign, though under a favorable sign additional capability is beneficial and maximization is not shown suboptimal; in the homogeneous pinned schedule used for the winner’s-irony result, the retained option forfeited by abandonment is nondecreasing in the gates already cleared; and the Golden Point supplies a viability constraint on alignment, not an objective.

What is imported. The survivorship companion [1] supplies the model, the two-condition theorem with its piecewise driver-viability functional, the corollary that adverse anatomy overlapping positive strength is fatal, and the ceiling. The technoanatomy companion [2] supplies the gating proposition in the pinned schedule model, the completion-weight fixed point, and the survival-time threshold. The Route World companion [3] supplies the Golden Point as the target and the Departure Point as its complement, the Disjointness Lemma and the Conditional IMD Corollary under the landscape hypothesis, and the residual ε of ε -Departure. Each is restated in Section 3 at its grade. The Nexic companion [4] is cited only for the historical-surprisal interpretation in the winner’s-irony section and for the graded external-hazard estimate; none of the core viability argument depends on it.

Position in the branch. The Golden branch is four papers in a strict order of import, and this is a fifth, downstream of all four: a consequence paper that owns very little machinery. It imports the survivorship, technoanatomy, and Route World companions and is cited by none of them; it does not use the theorem companion directly. The existential interpretation of the historical surprisal this paper separates from option value is developed in a later paper that cites this one; nothing here cites it, so that the citation graph, like the logical one, runs one way. Nothing in this paper is offered as support for the Route World hypothesis; where a conclusion depends on that hypothesis the dependence is stated.

What is refused. No certified favor: finite observation of a trajectory cannot certify that it is generated by the Golden-conditioned law, so nothing here claims the sign has been settled favorably or that any agent is aligned. No occupant: the argument that civilization is the biosphere’s only available catalyst is an argument about a role and not about who fills it, and no present is written as the decisive one. No agent behind the conditioning. And no derivation of the normative premise from the model; it is owned.

Grades. [P] proved here or in a companion; [S] scaling-level or structural, argued without proof; [Φ] a stipulative or evaluative commitment stated as such; [O] open, with its specification.

2 Vocabulary

- **Golden Point.** The state past which a biosphere’s external hazard is driven to zero and viability obtains; the target.
- **Departure Point.** Its complement: the irreversible loss of a biosphere’s future.

- **Technoanatomy.** The sign of what added capacity does to the internal hazard; the decisive variable.
- **Muffling agent.** The lineage whose capacity muffles the external channel; in the record, civilization.
- **Imperfection Magnitude Drop.** An alternative history whose joint probability with the target is drastically minute; under the landscape hypothesis, every alternative lineage.
- **ε -Departure.** Departure with a residual ε of recovery through a successor muffling agent; the price of the concession.
- **Environmental extinctionism.** Human extinction advocated because humans are held harmful to the biosphere, the biosphere’s persistence being the value in play; the paper’s target, and not anti-natalism as a position about procreation.
- **Capability-first techno-optimism.** The thesis that increasing capability is itself sufficient to solve the risks capability creates; the paper’s other target, and not every optimistic view of technology.
- **Premise Φ_{env} .** Greater long-run biospheric persistence is, ceteris paribus, environmentally preferable; the paper’s one normative premise, owned and not derived.
- **Viability boundary.** A preference-independent boundary supplied by the model; for an alignment objective constrained to preserve biospheric persistence, Premise Φ_{env} turns that boundary into a constraint, which is a necessary condition and not the whole objective.
- **The ceiling.** No event of bounded duration certifies which law a surviving record is under; hence no certification, from a trajectory, that an agent is generated by the Golden-conditioned law.
- **The successor scenario.** An artificial agent ending humanity and reaching the target for itself; priced at ε .

3 Imported results

Numbered premises. Where the text refers to $\Phi 2$, it is the framework’s scope declaration, stated in the Route World companion [3], that Golden reasoning is a vantage and not an obligation. This paper’s own normative premise is Φ_{env} alone; the meta-normative selection rule of the broader program is not part of its machinery.

Definition 3.1 (The model, the target, and its complement; from [1, 3]). The total hazard is $h_t = B_{t\zeta b}(x_t) + \Gamma e^{-\rho x_t}$, an external background muffled by accumulated capacity x plus an internal term the same capacity feeds or starves by the sign of ρ ; Departure is the killing time, viability the event that the total hazard ever accrued stays finite, and the Golden Point the state past which the external channel is driven to zero and viability obtains. The Route World companion constructs the Departure Point from the record’s side as the least fixed point of an implication cascade and shows the two constructions agree.

Theorem 3.2 (Two conditions for viability; [P] in [1]). *For $\rho \leq 0$ viability is null for every parameter value and every driver. For $\rho > 0$, $\mathbb{P}(\Lambda_\infty < \infty) = \mathbf{1}\{\rho > 0\} V_{\text{drv}}(\mu_0)$ with V_{drv} the piecewise driver-viability functional; the two factors are logically distinct necessary conditions,*

the first a sign and the second silent on it. No index of accumulated capacity distinguishes the ascent that arrives from the ascent that is a fall.

Corollary 3.3 (Adverse anatomy overlapping positive strength is fatal; [P] in [1]). *If for some $a, b > 0$ the set $\{t : \rho_t \leq -a, x_t \geq b\}$ has infinite Lebesgue measure almost surely, then $\Lambda_\infty = \infty$ almost surely.*

Proposition 3.4 (An unheld x leaves the crest at bare background; [P] in [2]). *In the pinned schedule model a biosphere that cannot hold x faces the external background unshielded and pays the internal channel's full coefficient, at total hazard $B_t + \Gamma$, and $\mathbb{P}(\Lambda_\infty < \infty) = 0$ identically, for every parameter value, every anatomy, and every driver.*

Proposition 3.5 (Completion weight as a least fixed point; [P] in the pinned schedule model, in [2]). *With $s = \lambda/(\lambda + B + \Gamma)$ the per-gate survival factor, the expected completion weight from one owed gate is the least root $u \in [0, 1]$ of $u = s(\pi + (1 - \pi)u^r)$, and from n gates it is u^n ; $u < 1$ whenever $B + \Gamma > 0$. Three probabilities are kept distinct here and named: $V = \mathbb{P}(\Lambda_\infty < \infty)$, that the underlying continuation has finite cumulative hazard; the probability of reaching such a continuation alive, which is what u^n scales; and $S = \mathbb{P}(T = \infty) = \mathbb{E}[e^{-\Lambda_\infty}]$, survival forever. A finite pre-completion hazard changes the second and not the first, so a product $u^n V_\infty$ is the probability of reaching a viable continuation alive and does not inherit V 's name. The floor $\Gamma > 0$ on the internal channel after capability is gone is a model assumption needing an interpretation independent of the maintaining civilization, taken here as external hazard that persists when the maintainer does not; where that interpretation is declined, the long-run result must be grounded in an explicit external-hazard assumption.*

Proposition 3.6 (The survival-time threshold; [P] given the model, in [2]). *Given survival to t , the posterior weight of histories that complete the schedule tends to one, past a threshold T_* logarithmic in the rarity of completion; conditioning on asymptotic survival lies past T_* for every parameter value and prefers the completing class outright, tilted toward better completion states.*

Lemma 3.7 (Disjointness; and the Conditional IMD Corollary, [S] under the landscape hypothesis, in [3]). *Every extant lineage other than our own is disjoint from the human route on its route attributes, and is therefore not the human route; under the additional hypothesis that sufficiently route-distant alternatives are Golden-suppressed, the cumulative-suppression package of the High-Contrast Earth Hypothesis, each such lineage's Golden history is an Imperfection Magnitude Drop relative to our own Aethus.*

Theorem 3.8 (The ceiling; [P] for the hard-wall law, [C] for the soft-wall law, in [1]). *In the critical regime the survivorship-conditioned law is locally absolutely continuous with respect to the unconditioned law on survivors and globally singular; no event of bounded duration perfectly separates the two, so finite behavioral observation cannot certify which generating law produced the record. What formal properties of an agent can separately be proved is a different question, and Section 5 keeps the two apart.*

4 The core result, and the successor as an optional strengthening

The argument against environmental extinctionism needs almost none of the speculative machinery this paper's companions develop, and it is stated first in the form that needs none. Remove

every available muffling capacity and the biosphere is the pinned case of Proposition 3.4: $x \equiv 0$, the external channel unshielded at B_t , the internal channel at its bare coefficient Γ , so that

$$h_t \geq \Gamma \implies \Lambda_\infty = \infty \implies \mathbb{P}(T = \infty) = 0,$$

for every parameter value, every anatomy, and every driver, with no assumption on the background's long-run behavior. That is the whole of the core. A reader who rejects everything downstream of the survivorship model can accept it.

The obvious objection is that another muffling agent might later appear, and the objection is met on its own ground rather than argued away. A successor route arriving at instantaneous rate α_t and completing the schedule with probability q_t on arrival competes against the pinned hazard $h^{(0)} = B + \Gamma$ that runs while it waits. With $r_t = \alpha_t q_t$ the probability that the first successful recovery beats the Departure clock is the competing-risks residual

$$\varepsilon_{\text{succ}} = \mathbb{E} \left[\int_{\tau}^{\infty} r_t \exp \left(- \int_{\tau}^t [h_u^{(0)} + r_u] du \right) dt \right], \quad (1)$$

which at constant rates is $r/(h^{(0)} + r)$ and never exceeds one. Deep time does not enter as a mechanism that alters a fixed perturbation weight; it enters through the competing-risks functional, a ratio of rates in the homogeneous case, and the claim that a post-extinction biosphere does not recover its route is the claim that r is small against $B + \Gamma$ along the path, for which the bound $r_t \leq \eta h_t^{(0)} \Rightarrow \varepsilon_{\text{succ}} \leq \eta/(1 + \eta)$ gives the time-varying form. The inputs α and q are unestimated, and the claim is entered at that grade. The Route World companion's stronger claim, that successor routes are Golden-suppressed under its landscape hypothesis, so that $\mathbb{P}(\mathcal{G} \mid \neg C)$ is drastically minute rather than merely reduced, is an optional strengthening carried at that hypothesis's grade (Lemma 3.7); the core result does not depend on it.

The perturbation the argument runs on is named exactly, because an earlier form of it was misdescribed. Extinction does not falsify any historical attribute of the record: after it, the lineage still descended through its ancestor and passed through its setting. What extinction removes is the future-option attribute

$$C = \{\text{the extant civilizational lineage remains available as a continuation}\},$$

and the perturbation is $C \mapsto \neg C$, with the quantity in question $\mathbb{P}(\mathcal{G} \mid \neg C)$ against $\mathbb{P}(\mathcal{G} \mid C)$. The historical facts, the last common ancestor and the climatic sequence that produced the lineage, enter only as the reasons C is not reproducible once lost.

5 Alignment: a viability constraint, not an objective

Artificial intelligence is a perturbation of the state (x, ρ) , and on the technoanatomy companion's design the largest perturbation of ρ in the record, delivered during the window in which the anatomy is being settled. That is the whole of what this paper says about it, and it says it as a constraint. The viability boundary of the two-by-two picture of Section 12 is preference-independent: it is a property of the model, and no party's preferences draw it. Whether an aligned system must preserve it is a separate question, and the answer is conditional in exactly the way the rest of this paper is: for any alignment objective constrained to preserve biospheric persistence, that is, under Premise Φ_{env} adopted for the system, the boundary is a constraint the system must not violate, by removing the muffling capacity or by turning the response sign

adverse, whatever else it optimizes. What the model contributes is the boundary, external to the alignment literature’s inventory of instructions, intentions, revealed and ideal preferences, interests, and values [5]; what makes it a constraint is the premise, and the paper does not let the mathematics carry that step. It determines nothing further. Two viable policies may differ arbitrarily in welfare, liberty, distribution, biodiversity, the treatment of individuals, and whose values prevail, and survival ranks none of those. Nor is the constraint certifiable from behavior: by Theorem 3.8, finite observation of a trajectory cannot certify that it is generated by the Golden-conditioned law, whatever a formal specification may separately be proved to satisfy. What the species does with the constraint, and what it would mean for a species to own its conduct in the light of the structure that produced it, is developed separately and is not this paper’s subject.

6 The concession priced: ε -Departure

Definition 6.1 (ε -Departure Point). A state is an ε -*Departure Point* when the total Aethic weight of continuations that advance is at most ε . A Departure Point in the strict sense of the Route World companion [3] is the case $\varepsilon = 0$.

7 The human role: two errors about it

Two positions occupy the ends of a spectrum concerning what humans are for. At one end, capability-first techno-optimism: capability is the thing, and more of it is the answer. At the other, a strain of environmentalism holding that humanity is a net harm to the biosphere and that its voluntary disappearance would serve the living world. Golden reasoning finds both mistaken, and mistaken in ways that are structurally opposite — each neglects exactly the factor the other over-weights.

The argument against techno-optimism is already in hand and is consolidated here: it neglects technoanatomy, and by the technoanatomy companion [2] capability growth without favorable anatomy cannot deliver viability and, once pushed beyond the adverse-anatomy hazard minimum, becomes actively self-defeating (Section 10). The argument against extinctionism is new and is the section’s substance: on the extinctionist’s *own* valuation, the biosphere’s persistence, removing the biosphere’s only hazard-lowering agent lowers the thing they mean to protect (Section 8). Both arguments are conditional, and the conditions are stated rather than assumed.

What emerges is a position on environmental value that neither end occupies: for the survivorship argument the valued object is the biosphere as a temporally extended living system, within the present survivorship model the Golden conjunction is the only route by which it acquires a positive probability of indefinite persistence, and human existence is on that accounting not an accident at odds with the living world but structurally implicated in its only presently realized route to positive viability (§11). The threat is not thereby diminished — $\rho < 0$ remains the live danger, and this section is as much an argument for urgency about anatomy as against despair about presence.

8 Against extinctionism

8.1 The position, in its own terms

The target is named narrowly. The position argued against is *environmental extinctionism*: human extinction advocated because humans are held to be harmful to the biosphere, the biosphere’s persistence being the value in play. Anti-natalism as a position about procreation, and suffering-based forms of it in particular, are outside the argument unless they separately endorse biospheric persistence as the relevant value; the argument runs on that value and on nothing else.

The target is a specific and explicitly biosphere-valuing position, and it must be stated in its own terms or the argument does not engage it. The Voluntary Human Extinction Movement holds that humanity should abstain from reproduction so as to bring about its own gradual disappearance, on the ground that this would prevent environmental degradation [6]. The reasoning proceeds *from* the value of the living world: the movement’s own framing weighs the greater value of Earth’s entire biosphere against the lesser value of one species, and holds that once one envisages the biosphere entire and grants the intrinsic value of every life form, voluntary extinction begins to make sense.

That framing is what makes the position addressable here. Its premise, that the biosphere’s persistence and flourishing is what matters, is one this framework can evaluate, because persistence of a biosphere is the quantity Definition 3.1 models.

What this argument does not touch. It is essential to distinguish the target from a neighboring position it is often assimilated to. Benatar’s antinatalism holds that coming into existence is always a serious harm, and concludes that human extinction would be preferable, on an asymmetry between the badness of pain and the goodness of pleasure [7]. That argument does not proceed from biosphere value at all; its currency is suffering, and extinction is a consequence rather than the objective.

Nothing below engages it. An argument that removing humans lowers the biosphere’s expectancy is no reply to someone who never valued the expectancy. The suffering-focused position must be met on its own ground and this section does not attempt it. What follows applies to extinctionism *in its environmental form* — and that restriction is not a weakening, since the environmental form is the one that appeals to a premise the framework shares.

8.2 The argument

Proposition 8.1 (Removing the muffling agent forfeits viability; [P] given the model). *In the model of Definition 3.1, a biosphere with no hazard-lowering agent is the pinned case $x \equiv 0$: no muffling is accumulated, the external channel is unshielded at B_t , and the internal channel sits at its bare coefficient Γ , so by Proposition 3.4 $h_t = B_t + \Gamma$. Then $\Lambda_t = \int_0^t (B_u + \Gamma) du \geq \Gamma t \rightarrow \infty$, with no assumption on the background’s long-run behavior needed, so*

$$\mathbb{P}(\Lambda_\infty < \infty) = 0$$

identically — not small, but zero, and for every parameter value. Human extinction is not literally the pinned case, and the bridge is stated as a lemma rather than assumed: automated infrastructure or an artificial successor could in principle hold x positive, and those branches belong to the successor analysis below and not to the no-capability branch. What extinction

of the maintaining civilization supplies, absent any autonomous replacement, is decay of the maintained capacity, $x_t \rightarrow 0$, and that suffices.

Lemma 8.2 (Capability decay; $[\mathbf{P}]$ given the model). *If $x_t \rightarrow 0$, then $\Gamma e^{-\rho x_t} \rightarrow \Gamma > 0$, so there is t_0 with $h_t \geq \Gamma/2$ for all $t \geq t_0$, whence $\Lambda_\infty \geq \int_{t_0}^\infty \Gamma/2 dt = \infty$ and $\mathbb{P}(T = \infty) = 0$.*

The premise the extinctionist position must therefore accept is that human extinction with no autonomous replacement makes $x_t \rightarrow 0$; if it denies that, it has moved to the successor branch and must price the replacement.

Compare the alternative. By Theorem 3.2, a biosphere that does accumulate muffling has $\mathbb{P}(\Lambda_\infty < \infty) = \mathbf{1}\{\rho > 0\} V_{\text{drv}}(\mu_0)$, which is positive whenever anatomy is positive and the driver-viability functional is: for a transient driver escaping toward $+\infty$, or a positive recurrent one with positive stationary mean, which widens the class of muffling dynamics this comparison can invoke. The extinctionist proposal is therefore a choice of an option with viability probability exactly zero over one whose viability probability may be positive. The word may carries a premise that is stated rather than left in the rhetoric: the mathematics proves that extinction forfeits the only modeled route, and strict preference under the persistence criterion requires in addition that the retained route have positive value,

Premise $\Phi 1$ (Retained option; $[\Phi]$). Write

$$V_{\text{current}} := \mathbb{P}_{S_{\text{now}}}(\text{enter the positive-viability region before Departure}),$$

the probability, from the present state S_{now} , that the retained civilizational route enters the live region of Theorem 3.2, the technoanatomy schedule’s surviving-completion kernel integrated against $\mathbf{1}\{\rho > 0\} V_{\text{drv}}$ supplying its exact form. The premise is $V_{\text{current}} > 0$: the anatomy is movable.

Without it the comparison is 0 against 0, and an extinctionist convinced that ρ is irreversibly adverse retains a coherent objection, which the section takes up rather than dismisses.

Two scopings keep the zero honest at the site. The $[\mathbf{P}]$ is model-internal: it prices the fixed-configuration model, and the model-to-world step remains this section’s ledger cost at $[\Phi]$. And the zero is configuration-conditional: the recovery branch, a successor lineage re-evolving the muffling capacity, is precisely the branch a fixed-biosphere model does not represent, and the corpus prices rather than ignores it: that branch is the ε of ε -Departure (Definition 6.1), where the pricing sentence lives.

The magnitude of what is forfeited can be given a number, and the number is graded at the level of the hypotheses that supply it and is not part of the core result. The companion treats the external hazard as roughly constant over long timescales, since the cosmic background does not change appreciably from era to era, giving a survival probability $e^{-\lambda t}$ and an external life expectancy of $1/\lambda$ [4]. Under the Rare Earth and High-Contrast Earth hypotheses, at their grade $[\mathbf{C}]$, the framework’s reading is $\lambda t_0 > 1$, where t_0 is the length natural history actually took, which would say that Earth’s remaining expectancy under the external hazard alone is shorter than the span already elapsed; the core argument above needs none of this.

Remark 8.3 (Undisturbed nature is not preserved nature). This is the point on which the position turns, and it is worth stating without ornament. The extinctionist picture is of a living world continuing undisturbed once the disturbance is removed. But “undisturbed” is not a stable state; in the closed pinned model it is the case $x \equiv 0$, strict Departure, in which the biosphere has

no means of lowering its own hazard and the background runs unopposed, and in the expanded model that admits successor recovery it is ε -Departure, the enlargement Remark 8.4 announces.

The reframing of the Route World companion [3] sharpens what is being chosen, and sharpens it against the position. Under a reading on which the relevant event is biological death, “leave nature alone” sounds like survival, and the argument must show that death follows anyway. Under the reading on which the Departure Point is the closure of possibility, no such demonstration is needed. At $x \equiv 0$, Proposition 8.1 gives $\mathbb{P}(\Lambda_\infty < \infty) = 0$ identically: *no* continuation available to that biosphere advances. The modal quantifier of the Route World companion [3] is therefore satisfied at once, and the Departure Point is not a fate awaiting the undisturbed world but its present condition. In the terms of the Route World companion [3]: $x \equiv 0$ lies in $\mu G \setminus \text{Base}$ — it is a departure at which no hazard has yet fired, entered by configuration rather than by event. The extinctionist proposal is a proposal to enter the cascade deliberately. A biosphere left alone has not been preserved; it has been placed in the one configuration from which nothing leads anywhere. The extinctionist is not choosing a world that survives without us. They are choosing a world that cannot go anywhere. Destabilization of the star, a sufficiently close hypernova, orbital perturbation by a passing mass — these do not become less likely because humanity has withdrawn. Nothing any longer counteracts them. the technoanatomy companion [2] reaches the same configuration from the other side: a biosphere that has not completed its schedule has not yet acquired what the extinctionist proposes to give up, and the two sit at the same bare background.

And the choice is not between a damaged biosphere and an undamaged one. It is between a biosphere with a positive probability of persisting indefinitely and one with none. On the extinctionist’s own valuation, the second is worse.

8.3 The concession, and why it does not rescue the position

The argument has a real limit, and stating it is what keeps it honest.

Remark 8.4 (Adverse anatomy does not vindicate the position). The natural concession at this point is that under $\rho < 0$ the extinctionist is closer to right, since $x \equiv 0$ leaves hazard merely at background while adverse anatomy drives it upward. The concession should not be made, and seeing why changes which options the argument is really between.

First, the local comparison goes the other way. A biosphere with adverse anatomy that *halts* at the optimum x_* of the survivorship companion [1](iii) sits at hazard h_* , and $h_* \leq h(0)$ for every $\rho < 0$, strictly except at the tuned parameter values at which $x = 0$ is itself the minimum. Numerically, at $B = 1$, $b = 3$, $\Gamma = e^{-6}$:

ρ	$h(0)$	x_*	$h(x_*)$	expectancy gain
−0.05	1.00	11.47	0.0046	217×
−0.20	1.00	8.88	0.0176	57×
−1.00	1.00	4.28	0.407	2.5×
−2.00	1.00	1.82	0.898	1.1×

Even at $\rho = -2$ — anatomy adverse enough that capability doubles internal hazard twice over — halting still beats having no muffling at all. So $x \equiv 0$ is weakly dominated for every $\rho < 0$, and strictly dominated generically. Adverse anatomy is worse than absence only if the biosphere *fails to halt*, which is a failure to stop rather than a property of the anatomy.

Second, and this is the point the local comparison obscures, both options forfeit. By Theorem 3.2, $\mathbb{P}(\text{viable}) = \mathbf{1}\{\rho > 0\} V_{\text{drv}}(\mu_0)$. At $x \equiv 0$ there is no muffling agent, hazard sits at

background, and $\Lambda_\infty = \infty$. At $\rho < 0$ halted at x_* , hazard sits at $h_* > 0$, and $\Lambda_\infty = \infty$. *Both are zero.* The extinctionist and the biosphere that settles at x_* are not choosing between a worse and a better outcome; they are choosing between two ways of dying, and the model does not distinguish them in the only respect that admits an answer other than death.

The response to adverse anatomy is therefore not to select among the zeros. It is to move ρ , which is the one intervention that changes the indicator rather than the multiplicand — and which the broader Golden program exists to model. A reader persuaded of the danger and unpersuaded that ρ can be moved still has a coherent position, and it remains the strongest form of the view opposed here; but it is a claim about the tractability of anatomy, not a defense of extinction.

Third, the recovery branch is modeled rather than named. Elapsed time does not by itself restore a route, but neither is an Imperfection Magnitude Drop invariant under arbitrarily many future opportunities: a successor muffling agent arriving at a small instantaneous rate has, over long enough intervals, a cumulative probability that is not small. Let α_t be the arrival intensity of a candidate successor and q_t its probability of completing the schedule on arrival, and let $h^{(0)}$ be the pinned hazard $B + \Gamma$ of the interval before recovery. In the competing-risks form the residual of Definition 6.1 is

$$\varepsilon_{\text{succ}} = \mathbb{E} \left[\int_{\tau}^{\infty} r_t \exp \left(- \int_{\tau}^t [h_u^{(0)} + r_u] du \right) dt \right], \quad r_t := \alpha_t q_t, \quad (2)$$

the probability that the first successful recovery beats the Departure clock, the two hazards competing independently; for constant rates it is $r/(h^{(0)} + r)$, which the small- r approximation $r/h^{(0)}$ reproduces and which never exceeds one. The claim that deep time does not rescue the position is then the claim that (2) is small, and for time-varying rates it depends on the whole intensity path; a bound that does not assume constant rates is available: if $r_t \leq \eta h_t^{(0)}$ pointwise, then the recovery cause occupies at most the fraction $\eta/(1 + \eta)$ of the competing hazard, so that, provided the competing total hazard eventually fires, $\varepsilon_{\text{succ}} \leq \eta/(1 + \eta)$. Recovery must outrun Departure, and that is a statement about the relative intensity profile, reducing to a ratio of rates in the homogeneous case, not about the passage of time; it is entered at that grade, with α and q unestimated [O]. One enlargement of the model is announced here because it changes a name: in the closed pinned model the state $x \equiv 0$ is strict Departure, no advancing continuation existing; once successor recovery is admitted with $\varepsilon_{\text{succ}} > 0$ the same configuration is an ε -Departure state in the sense of Definition 6.1, and the two scopes are not run together.

8.4 Why forfeiting is not the same as never having had

Purpose and its grounding are treated once, at the broader Golden program; this subsection confines itself to stakes. The preceding leaves a gap that arithmetic alone will not close. If both options carry $\mathbb{P}(\text{viable}) = 0$, why prefer either to the other, and why regard the choice as weighty rather than indifferent? The answer is that the two are not symmetric with respect to what has already been spent — and note that under the Route World companion [3] what is at stake is the trajectory rather than the lives on it, which is why the accumulation is the right object to weigh. A biosphere at $x \equiv 0$ has not merely a low viability probability; it has *surrendered* an accumulation, and the accumulation is the subject of the companion's Miracle sequence [4]. Whatever its correct magnitude, the record is a chain of events each of which had to succeed for the next to be attempted: an origination, a transfer if panspermia obtained,

the eukaryotic transition, the Dark Era, the end-Cretaceous impact, the α -gal epidemic. The companion's estimate for a single such step, the interstellar transfer, is of order 10^{-21} [4].

Remark 8.5 (Forfeiture is not observable). The cascade of the Route World companion [3] adds a feature to the stakes argument that deserves separate statement, because it is the feature that makes the argument urgent rather than merely large. A biosphere enters μG at the moment the last advancing route closes, and by the Route World companion [3] that moment precedes the firing of any hazard — possibly by a great deal. Nothing marks it. There is no event to observe, no loss to register, and by Theorem 3.8 no event of bounded duration that would certify that a biosphere has just crossed rather than not. The accumulation is not forfeited in some future catastrophe one might see coming and reconsider; it is forfeited silently, at whatever point the routes ran out, and the catastrophe is only the execution.

This does not make the stakes argument more permissive, see the limits below, but it does relocate what the argument is about. It is not a claim that a visible disaster would be very bad. It is a claim that the relevant transition is invisible, which is why deferring attention to anatomy until the danger is legible is deferring it past the point where deferral is possible.

Remark 8.6 (What this does and does not license). The stakes argument establishes an asymmetry between forfeiting and never having had, and nothing beyond it. In particular it does not license: that any present cost is acceptable, since the argument is silent on rates of exchange; that the Golden Point will be reached, which Theorem 3.8 forbids anyone from knowing; or that persons alive now bear a duty derived from the accumulation, which $\Phi 2$ excludes and which the framework has no machinery to generate.

What it does license is narrow and, we think, sufficient: that the two options which forfeit viability are not the neutral defaults they are often taken for, and that between an option which changes the indicator and options which do not, the model has a determinate preference.

9 The winner's irony: the stakes of forfeiture

Proposition 9.1 (The stakes of forfeiture rise with the gates cleared; [P] in the homogeneous pinned schedule model with a common completion state). *Let the schedule be homogeneous and pinned, with completion weight $u < 1$ per gate (the technoanatomy companion [2]), and suppose every completing branch enters a common deterministic completion state of viability $V_\infty = \mathbf{1}\{\rho > 0\}V_{\text{drv}}$; without that idealization $u^{n-k}V_\infty$ is replaced by the integral of V_{drv} against the surviving-completion kernel from stage k , and the monotonicity below is the kernel's. For a lineage that has cleared k gates, the viability forfeited by abandoning the schedule is*

$$V_k = u^{n-k} V_\infty - 0, \quad (3)$$

since abandonment is the pinned case with viability zero; V_k is nondecreasing in k . The unconditional probability of having reached stage k is $\prod_{i \leq k} p_i$, nonincreasing in k , and its surprisal $\sum_{i \leq k} (-\log p_i)$ is nondecreasing. So along the trajectory two quantities rise together, and they are of different kinds: what is held, V_k , which is the stakes, and how unlikely it was to be held, $\sum_{i \leq k} (-\log p_i)$, which is historical surprisal. Only the first establishes stakes; the second is the separate fact whose existential reading belongs to the Miracleism material of the broader program, and neither is inferred from the other.

Proof. Abandonment at stage k is Proposition 3.4's pinned case, viability zero; continuation from stage k has $n - k$ gates remaining, each costing the factor u , and then V_∞ . Monotonicity in k is $u < 1$. The probability and surprisal statements are definitional. $\square$

Two consequences follow, and one correction of an earlier form of the argument. Within this homogeneous gate schedule, the stakes attaching to the decision to continue rise along the trajectory, so a later stage carries higher stakes than an earlier stage, not by self-importance but by (3); this monotonicity is a result of the schedule model, not a universal theorem that every later historical moment must carry greater stakes than every earlier one; what does not rise is the probability of failing at a given step, and an earlier version of this remark that read the product $\prod_{i < n} p_i$ as an increasing cost had the direction backwards, since the product falls. And sunk cost is not the structure here: the classical sunk-cost fallacy consists in letting past expenditure govern a choice whose future payoffs are unchanged by it, whereas here the past expenditure has *purchased* the option V_k that remains, and forfeiting is forfeiting that option rather than honoring a debt. The distinction is that between a bad reason to persist and a live asset one may discard.

Remark 9.2 (The winner’s irony). Proposition 9.1 gives the irony in desiring the end of civilization for nature’s sake its exact form. On geological timescales the lineage has only just cleared the gates that make the target reachable, an outcome the record prices at the surprisal $\sum_i (-\log p_i)$ of every gate behind it, the scale of which the Nexic companion’s own accounting of the Cenozoic chain conveys [4]; and what that improbability put in hand is the option V_k , the biosphere’s only realized route to positive viability. To wish the route dismantled for the biosphere’s sake is to hold a winning ticket and push the winnings away out of ignorance of what has just happened, with the difference that the ticket is not inconsequential: the stakes are the continuation of the living world, and, within this homogeneous schedule model, they rise with every gate cleared, which is why the proposition and not sentiment carries the point.

10 Against techno-optimism

The opposite error requires less new argument, since the result it founders on is already proved. Techno-optimism holds that continued growth in capability is the answer to the risks capability creates. In the model’s terms it supplies the driver-viability factor of Theorem 3.2 and is silent on the anatomy factor — ρ does not appear in any accounting of energy capture, capability, or rate of progress, and cannot be inferred from them.

The verdict is not that this leaves the conclusion unsupported but that it inverts it. With $\rho < 0$, feedback does not fail to secure survival; it *accelerates the failure*, since $\mu_t \rightarrow +\infty$ drives $x_t \rightarrow +\infty$ and hence, by the survivorship companion [1](iii), $h \rightarrow \infty$ superexponentially. Capability alone cannot deliver positive viability under adverse anatomy, and beyond the adverse-anatomy hazard minimum x_* continued growth becomes actively self-defeating: a poorly-constituted biosphere that halts sits at its hazard floor $h_* \leq h(0)$, while one that accelerates past x_* runs at its own cliff. Both have viability zero; what separates them is the instantaneous hazard along the way.

Remark 10.1 (The cascade forces anatomy to be a process). Both arguments of this section, taken with the Route World companion [3], deliver a verdict on the fixed- ρ model that is worth confronting rather than absorbing. If $\rho < 0$ and ρ is a *constant*, then Theorem 3.2 gives $\mathbb{P}(\Lambda_\infty < \infty) = 0$ identically — at every time, including $t = 0$. By the modal definition the biosphere is then in μG from the outset: no continuation available to it ever advances, and it departed before anything happened. That is a degenerate verdict. It says the question of what such a biosphere should do was settled before it was asked, which is not a claim about biospheres but a symptom of the model.

The degeneracy has exactly one exit, and it is not a matter of taste. It is avoided if and only if ρ can change — if adverse anatomy now leaves open a continuation in which anatomy is not adverse later. That is the fluctuating-anatomy setting of the survivorship companion [1], where Corollary 3.3 and the survivorship companion [1] take over, and where the question becomes whether the escape is completed in time rather than whether it was ever available.

So the cascade converts emergent anatomy from a desirable extension into a *requirement of coherence*. A framework that treats technoanatomy as a parameter and departure as modal closure will pronounce every adverse-anatomy biosphere already lost; only a framework in which anatomy has its own dynamics can say the thing this section is trying to say, which is that the matter is open and the window is closing. The arguments above should be read as holding under the fluctuating reading, and the fixed- ρ statements as the limiting case in which the window has already shut.

The two errors are the same error. Placed side by side, the positions are not opposites but a single mistake taken in two directions. Both treat capability as the only variable. The techno-optimist maximizes it; the extinctionist minimizes it; neither carries a parameter for whether capability, once acquired, is turned inward. Once ρ is in the model the spectrum they occupy is revealed as one axis of a two-dimensional space, and the interesting region, positive anatomy with accumulating capability, lies on neither end of it. That is why the disagreement between them has proved unresolvable by argument along their shared axis: the quantity that settles it is not on that axis.

The refutation’s arsenal has grown since this subsection was written, and the additions belong on display here. The invariant of the technoanatomy companion [2] says the operative threshold does not live in the space capability measures at all; the grabby engagement says visibility-weighted selection is not persistence-weighted selection, so the loud are not thereby the lasting (the technoanatomy companion [2]); and $\beta < \alpha$ is the one inequality on which capability-first hope depends and which its proponents never argue (the technoanatomy companion [2]). One further point is sharper than an arsenal item and is entered with its edge showing. As an interpretive observation and not a part of the refutation, the escalator-myth checklist this framework was run against (the broader Golden program) matches capability-first techno-optimism at several points, guaranteed ascent, salvation through capability, the future as reward, the structure Midgley named [8]; the two-condition theorem is the specific fact that structure does not check. The observation adds nothing to the theorem and is entered as interpretation.

11 What this implies for environmental value

Premise $\Phi 2$ (Environmental value; $[\Phi]$). *Ceteris paribus*, greater long-run persistence of the biosphere, considered as a temporally extended living system, is environmentally preferable. Everything this subsection concludes is conditional on this premise and on the model; the premise is the extinctionist’s own valuation stated as a premise, and nothing here derives it from the model.

The premise names its object narrowly. For the purpose of the survivorship argument the valued object is the biosphere as a temporally extended living system, and no identification of “nature” with it is claimed: abiotic systems, wilderness, ecological integrity, species diversity, processes independent of design, individual organisms, and sentient welfare are each things environmental

thought has meant by the word, and the result below is about biospheric persistence and not a total theory of environmental value.

Three consequences follow, of decreasing formality and increasing interest. They are stated as consequences of the foregoing and inherit its conditions. *The biosphere is the valued object of this argument.* The unit the Route World companion [3] adopts for methodological reasons, the totality of living and technical activity associated with a world, taken as one system, is what some holistic environmental traditions treat as one important object of environmental value, the whole rather than particular organisms; no claim is made that it exhausts what environmental thought means by nature, and Premise $\Phi 2$ names the object narrowly. The framework arrives at the same unit from the requirement that a conditioning target apply across the whole record (the Route World companion [3]).

Positive viability requires the Golden conjunction. In the model, everything outside the live region is nonviable. A biosphere without a hazard-lowering agent is the pinned case, with the hazard floor $B + \Gamma$; a biosphere with adverse anatomy is fatal with certainty by Theorem 3.2. Only the conjunction of positive anatomy, favorable driver dynamics, and accumulating capability yields $\mathbb{P}(\Lambda_\infty < \infty) > 0$, and many trajectories could in principle instantiate that region; what the Route World companion adds is the identification of its modal target as the Golden Point. The criterion is positive probability of indefinite persistence, not infinite expected lifetime, which the critical regime has without viability. Whatever else environmental value recommends, if it recommends the living world’s persistence then it recommends passage through that conjunction, since no other route is available within the model. The modal form of the stronger claim, that passage through the Golden Point is the only route by which the Departure Point is avoided, is the broader Golden program’s, where it rests on the complementarity of the Route World companion [3] rather than on the expectancy comparison made here.

Human existence is not ontologically at odds with the living world. This is the claim most likely to be over-read, so its content and its limits are given precisely.

Remark 11.1 (What is and is not established). *Established, conditional on the foregoing:* the outcome in which the biosphere persists indefinitely runs through the existence of an agent capable of lowering its hazard. Humanity is, so far as the record shows, that agent’s occurrence in this biosphere. So on this accounting human existence is not contingent noise afflicting an otherwise self-sufficient nature; it is structurally implicated in the only presently realized route by which the biosphere acquires positive viability. Under Golden-conditioning (the Route World companion [3]) the implication is stronger still, since the record is then read as produced under a conditioning that the passage occur.

Not established: that the harm is small, that any particular human activity is vindicated, or that the present trajectory is good. The threat is real, $\rho < 0$ is the live danger, and Remark 8.4 declines the concession that a biosphere with adverse anatomy is worse than none, and the refusal is an argument rather than an optimism. The claim concerns humanity’s *structural position*, not its conduct, and the two are independent: an agent can be necessary to an outcome and be failing at it.

Also not established, and worth disclaiming explicitly: any conclusion about the standing of human meaning-making at large. It is tempting to read the above as supplying, from a source outside human self-interpretation, a verdict that humanity is not a mistake — and as answering a broader cultural pessimism thereby. The temptation should be resisted at this grade. What the framework delivers is a conditional structural claim about one biosphere’s route to persistence, resting on the framework’s biospheric-referent premise, stated in the Route World companion

[3], and on the Golden-conditioning argument, both of which are argued rather than proved. That is a substantial thing to have. It is not a philosophy of history, and presenting it as one would trade a defensible result for an indefensible one. The existential interpretation, developed separately, makes a claim in this neighborhood and does not disturb the refusal: what it offers is conditional on an adopted premise rather than issued by the apparatus, and the difference between a conditional grounding and a verdict from outside is the distinction this remark is drawing.

12 The two-by-two picture

The two results assemble into one table, which is the paper’s core and the reason neither end of the spectrum is the model’s optimum.

	$\rho \leq 0$	$\rho > 0$
capability unavailable ($x \equiv 0$)	nonviable	nonviable
capability accumulating	nonviable	potentially viable

The first row is Proposition 3.4: with capability pinned, the hazard floor $B_t + \Gamma$ makes viability zero at either sign, which is the result against environmental extinctionism. The second row’s first cell is Theorem 3.2 at $\rho \leq 0$, viability null for every driver, which is the result against capability-first techno-optimism. The remaining cell is the only one with positive viability, and it is positive only when the driver-viability functional is: the model’s live region is accumulating capability with favorable anatomy, and the two errors are the two ways of leaving it.

13 What is not claimed

Three fences hold everywhere above and are stated once. No certified favor: by Theorem 3.8 finite observation of a trajectory cannot certify that it is generated by the Golden-conditioned law, so nothing in this paper establishes that the response sign has been settled favorably, that civilization will reach the Golden Point, or that any agent, human or artificial, is aligned to it in that sense; every favorable conclusion is held and none is possessed. No occupant: nobody is written as the Golden Point’s arrival or its agent, the present is not written as the decisive moment, and the argument that civilization is the biosphere’s only available catalyst is an argument about a role, not about who fills it. No agent behind the conditioning: the conditioning is a statement about which histories are in the support and with what weight, and the paper invokes nothing that acts on the past. Where a conclusion depends on the Route World hypothesis, that our record is drawn under conditioning on the target, the dependence is stated at the site and the conclusion carries the hypothesis’s grade.

14 Open problems

Every open item is gathered here. The landscape hypothesis under which every alternative lineage is an Imperfection Magnitude Drop is the Route World companion’s and is open there. The successor residual $\varepsilon_{\text{succ}}$ is modeled as a competing-risks integral and its inputs, the arrival intensity and completion probability of a candidate successor, are unestimated; the claim that deep time does not rescue the extinctionist position is the claim that the integral is small,

and it is entered at that grade. The option-value proposition is proved in the pinned schedule model and is a surrogate for the leaky refinement. Whether alignment as a technoanatomy question admits an audit, a measurement of what each added capacity does to the sign, is the technoanatomy companion’s open identification map. And the Route World hypothesis itself, on which the strongest form of the anti-extinctionist conclusion depends, is a hypothesis with a prospective test and has not been tested.

15 Falsifiers

What would count against the paper’s claims, in one place. Against the anti-extinctionist argument: a configuration in which removing the muffling agent leaves the biosphere’s viability positive, which Proposition 3.4 forbids in the pinned model and which a demonstrated alternative muffling agent, one that clears the schedule without civilization, would supply; or an alternative lineage shown to reach the space-faring capacity by a different route, which would break the Conditional IMD Corollary’s premise. Against the anti-techno-optimist argument: a configuration in which capacity alone, with adverse sign, delivers viability, which Theorem 3.2 forbids. Against the alignment constraint: evidence that the claimed viability boundary itself depends on an unstated, preference-dependent modeling choice; against the certification claim, a finite trajectory shown to separate the two laws with probability one, which the ceiling forbids. Against the successor pricing: an estimate of ε that is not small, which would make the recovery branch a live route rather than a residual.

16 Conclusion

The claim stated at the outset stands at its grades, and it is a conditional. Under Premise Φ_{env} and the model, removing the mechanism that maintains the biosphere’s hazard-muffling capability, so that the capability decays to zero, cannot be an environmental improvement, because a decayed capability leaves a permanent hazard floor and viability exactly zero; maximizing capability cannot realize the objective either, because no magnitude of capacity substitutes for a favorable response sign; the two errors are the two ways of leaving the model’s one live region; and, in the homogeneous pinned schedule used for the winner’s-irony proposition, the continuation option forfeited by abandoning the realized route is nondecreasing in the gates already cleared. The concession is priced rather than argued away, and its price is a competing-risks functional, reducing to a ratio of rates in the homogeneous case. For alignment the model supplies a preference-independent viability boundary; treating that boundary as a normative constraint depends on the paper’s explicit persistence-valuing premise, and the boundary is held and never certified. What the model establishes about the question in the title is therefore this: civilizational disappearance is not the restoration of a self-sustaining biosphere. It is the removal of the currently realized route by which exogenous hazard can be endogenously reduced, a route whose cleared gates carry large historical surprisal without that surprisal itself constituting the forward-looking stakes, and the living world is the stake.

References

- [1] Ajax Benander. *Survivorship Conditioning for Hazard Rates Driven at the Log-Derivative Level*. Working draft. The Golden branch’s root extraction; the two-condition viability theorem with its

- piecewise driver-viability functional, the hard-wall and soft-wall conditioned laws, the certification ceiling, and the half-time theorem. 2026. URL: <https://aethus.org>.
- [2] Ajax Benander. *Technoanatomy: Civilizational Survival as a Sign, Not a Size*. Working draft. The mechanism extraction; the prerequisite schedule, the surviving-completion kernel, the survival-time threshold, and the front-loaded settlement result. 2026. URL: <https://aethus.org>.
 - [3] Ajax Benander. *The Route World: A Golden-Reasoning Hypothesis of Biospheric Natural History*. Working draft. The terminal synthesis of the technical branch; the Golden Point, the Departure Point, the Continuity Principle, and the conditioned-record hypothesis with its prospective test. 2026. URL: <https://aethus.org>.
 - [4] Ajax Benander. “Nexic Reasoning: Defining a Generalized Calculus Over Anthropic Parameters”. In: (Dec. 2024). DOI: [10.31219/osf.io/jvkcp](https://doi.org/10.31219/osf.io/jvkcp). URL: osf.io/jvkcp.
 - [5] Iason Gabriel. “Artificial intelligence, values, and alignment”. In: *Minds and Machines* 30 (2020), pp. 411–437.
 - [6] Les U. Knight. *The Voluntary Human Extinction Movement*. Movement founded 1991. 1991. URL: <https://vhemt.org>.
 - [7] David Benatar. *Better Never to Have Been: The Harm of Coming into Existence*. Oxford University Press, 2006.
 - [8] Mary Midgley. *Evolution as a Religion: Strange Hopes and Stranger Fears*. London: Methuen, 1985.